

Multimedia Retrieval

Practical Assignment

Irina Kiefer, Ying-Kai Dang


1. July 2020

Encoder

- `leaky_relu(Conv2d(3, 13, kernel_size=(1, 1), stride=(1, 1)))`
- `leaky_relu(Conv2d(13, 26, kernel_size=(2, 2), stride=(2, 2)))`
- `leaky_relu(Conv2d(26, 38, kernel_size=(2, 2), stride=(2, 2)))`
- `leaky_relu(Conv2d(38, 45, kernel_size=(2, 2), stride=(3, 3)))`
- `flatten`
- `leaky_relu(Linear(in_features=405, out_features=10, bias=True))`


Decoder

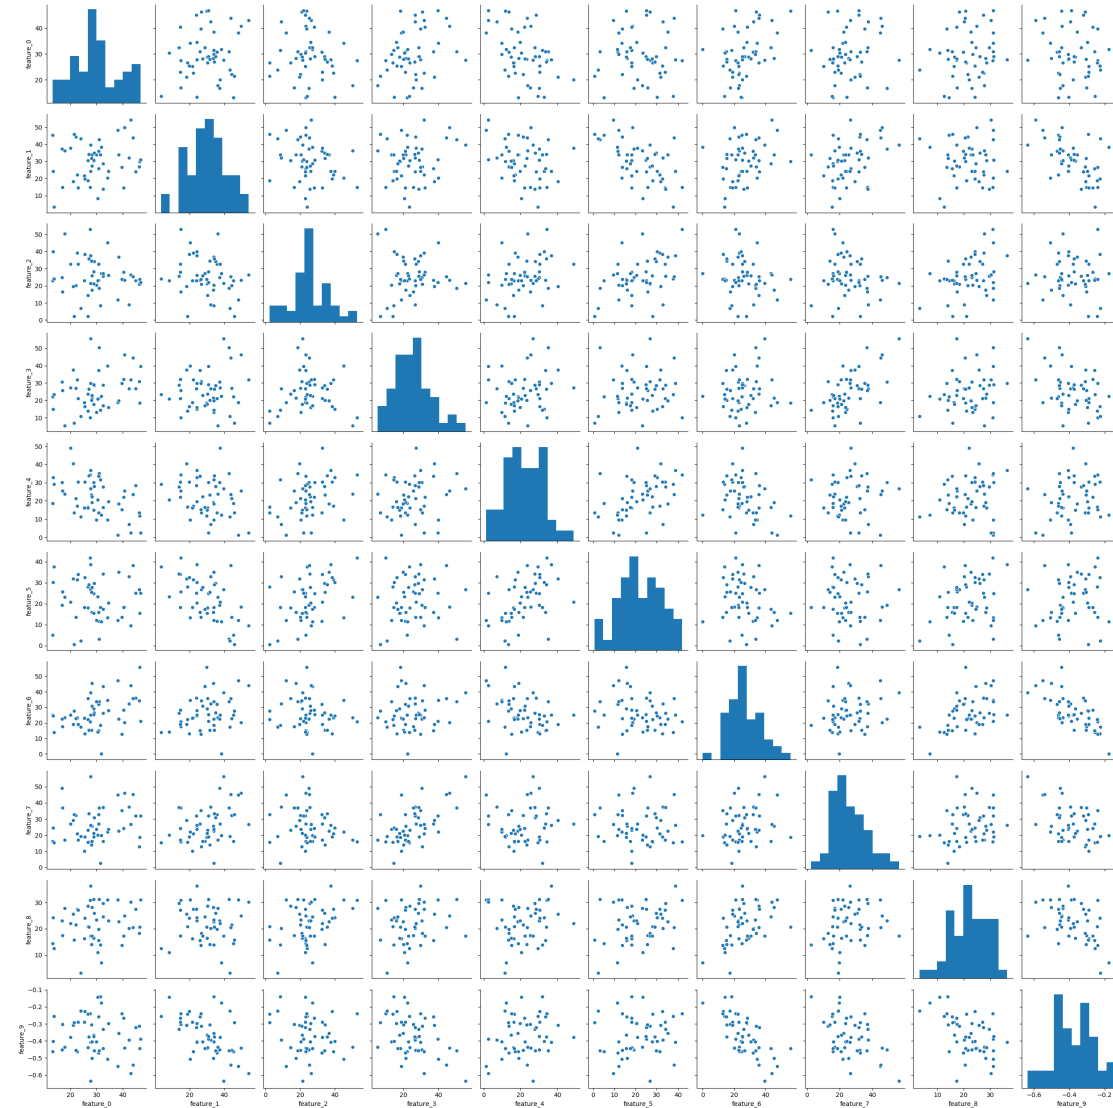

- `leaky_relu(Linear(in_features=10, out_features=16, bias=True))`
- `reshape((1, 1, 4, 4))`
- `leaky_relu(ConvTranspose2d(1, 5, kernel_size=(2, 2), stride=(1, 1)))`
- `leaky_relu(ConvTranspose2d(5, 9, kernel_size=(2, 2), stride=(2, 2)))`
- `leaky_relu(ConvTranspose2d(9, 14, kernel_size=(3, 3), stride=(2, 2)))`
- `leaky_relu(ConvTranspose2d(14, 19, kernel_size=(2, 2), stride=(1, 1)))`
- `flatten`
- `sigmoid(Linear(in_features=9196, out_features=3072, bias=True))`
- `reshape((1, 3, 32, 32))`



Trained on 50000 images over 2 epochs

Performance





UMAP & SPLoM